

Natural Language Processing: MT & Word Alignment Models



Christopher Manning

Borrows some slides from Kevin Knight, Dan Klein,
and Bill MacCartney



Lecture Plan

1. A bit more course overview [5 mins]
2. Briefly learn the history of Machine Translation [10 mins]
3. Learn about translation with one example [10 mins]
4. Speech recognition & Applying it to MT: The noisy channel model [10 mins]
[Stretch! **Emergency time reserves:** 5 mins]
5. Parallel-text word alignments: the IBM models [30 mins]



Lecture Plan

1. **A bit more course overview [5 mins]**
2. Briefly learn the history of Machine Translation [10 mins]
3. Learn about translation with one example [10 mins]
4. Speech recognition & Applying it to MT: The noisy channel model [10 mins]
[Stretch! **Emergency time reserves:** 5 mins]
5. Parallel-text word alignments: the IBM models [30 mins]



Learning Goals

- 1. To be able to **write** a natural language processing system which isn't just a keyword spotter or a bag-of-words text classifier**
2. To have a good sense of the complexity of natural language understanding (by either computers or humans)
3. To be aware of different levels of linguistic representation: words, syntax, semantics and tools used to model them
4. To be familiar with major approaches to natural language processing: rule-based, statistical NLP, other machine learning approaches including deep learning
5. To think that human languages and NLP are cool 😊



The class

Topics

- The course is not all MT! But I find it a fun place to start...
- Syntactic parsing, coreference resolution, named entity recognition, computational semantics, applications
- Classifiers, probabilistic models, deep learning, sequence models, generative and discriminative models applied to NLP
- Organization
 - Make sure you're on OpenEdX, Piazza, mailing list, website
 - Read about grading, collaboration, honor code, etc.
 - Programming assignment 1 (MT) is out today (!)
 - Need to use Java (except Final Project); encouraged to work in pairs

Do some reading!

- Jurafsky and Martin, Chapter 25: MT
 - Great for big picture of MT, okay for Model 1, no Model 2
- Adam Lopez, Statistical Machine Translation
 - Fairly modern, comprehensive Stat MT survey
- ~~Brown et al. 1993 “The Mathematics of Statistical Machine Translation”~~
- Kevin Knight, A Statistical MT Tutorial Workbook
 - Professor’s old favorite, but maybe time to move on? No model 2
- Michael Collins, Statistical MT: IBM Models 1 & 2
 - Closest to assignment/lecture notation. Maybe start here?
- Philip Koehn, *Statistical MT*, ch. 4: Word-based Models
 - Makes it too easy?!? Gee, there’s pseudo-code there...



Lecture Plan

1. A bit more course overview [5 mins]
2. **Briefly learn the history of Machine Translation [10 mins]**
3. Learn about translation with one example [10 mins]
4. Speech recognition & Applying it to MT: The noisy channel model [10 mins]
[Stretch! **Emergency time reserves:** 5 mins]
5. Parallel-text word alignments: the IBM models [30 mins]



MT: Just a Code?


“Also knowing nothing official about, but having guessed and inferred considerable about, the powerful new mechanized methods in cryptography—methods which I believe succeed even when one does not know what language has been coded—one naturally wonders if the problem of translation could conceivably be treated as a problem in cryptography. When I look at an article in Russian, I say: ‘This is really written in English, but it has been coded in some strange symbols. I will now proceed to decode.’ ”

Warren Weaver (1955:18, quoting a letter he wrote in 1947)





“When I look at an article in Russian, I say: ‘This is really written in English, but it has been coded in some strange symbols. I will now proceed to decode.’ ” – Warren Weaver, March 1947



“... as to the problem of mechanical translation, I frankly am afraid that the [semantic] boundaries of words in different languages are too vague ... to make any quasi-mechanical translation scheme very hopeful.”

– Norbert Wiener, April 1947



The early history: 1950s

- Early power
- Found proba
- First s
- MT he subst
- Little sema
- Proble





Machine Translation History

[see <http://www.hutchinsweb.me.uk/history.htm>]

1950s: Intensive research activity in MT

1960s: Direct word-for-word replacement

1966 (ALPAC): NRC Report on MT

- Conclusion: MT not worthy of serious scientific investigation; focus on understanding language

1966–1975: MT winter

1975–1985: Resurgence in Europe and Japan

- Domain-specific rule-based systems

1985–1993: Gradual Resurgence in the US

1993–2012: Statistical MT surges! Field becomes popular

2013: Out of data, might need new theory/models?

2014–2015: People excited about new neural MT models



What has happened between ALPAC and Now?

- Need for MT and other NLP applications confirmed
- Computers have become faster, more powerful
- Hugely increased availability of data: WWW
- People understand more about linguistics
- Change in expectations of quality
- Development of empirical, data-intensive, statistical, hybrid statistical/grammar-based, and neural approaches



Lecture Plan

1. A bit more course overview [5 mins]
2. Briefly learn the history of Machine Translation [10 mins]
3. **Learn about translation with one example [10 mins]**
4. Speech recognition & Applying it to MT: The noisy channel model [10 mins]
[Stretch! **Emergency time reserves:** 5 mins]
5. Parallel-text word alignments: the IBM models [30 mins]

الأخبار - عربي - حكومة مصر تقر تعديل قانون الانتخاب

http://www.aljazeera.net/NR/exeres/5FC61853-D3A2-4E68-915E-1014856CE98F

الجزيرة نت | الأخبار | الشخصية | المعرفة | الاقتصاد والأعمال | حريات وحقوق | مواقع أخرى للجزيرة

ابحث عن |

تواصل | مركز خريطة | بيانات | إلى |

العضوية | من |

شاهد البث الحي | لقناة الجزيرة

البث الحي

الأخبار

الاثنين 28/10/1432 هـ - الموافق 26/9/2011 م (آخر تحديث) الساعة 1:48 (مكة المكرمة)، 22:48 (غرينتش)

الصفحة الرئيسية: عربي

القائمة الرئيسية

عربي

تفاعلات

خدمات الموقع

حكومة مصر تقر تعديل قانون الانتخاب

شبان مصريون يرفعون لافتة تصف القانون الانتخابي قبل تعديله بأنه مضلل (الجزيرة نت)

أقرت الحكومة المصرية الأحد تعديل القانون الانتخابي الذي على أساسه ستجرى الانتخابات البرلمانية المرتقبة قبل نهاية العام، وغلبت فيه الاقتراع بنظام القائمة على النظام الفردي، وذلك بعد اعتراضات قوية على النسخة السابقة من القانون.

ونص القانون المعدل على انتخاب ثلثي أعضاء مجلسي الشعب والشورى بنظام القائمة الحزبية المغلقة، والثلث المتبقي بالنظام الفردي، وأن تتضمن كل قائمة مترشحة سيدة واحدة على الأقل.

وكانت النسخة السابقة من القانون -التي ردها المجلس العسكري الأعلى السبت إلى مجلس الوزراء- تقضي بانتخاب نصف أعضاء مجلسي الشعب والشورى بنظام القائمة الحزبية، والنصف الآخر بالنظام الفردي الذي تقول أحزاب إنه يسمح لفلول **الحزب الوطني الديمقراطي** المنحل بالعودة.

شاهد البث الحي

محتويات

تعديل مقترح لنظام انتخابات مصر

اقتراع مصر بنوفمبر ولقاءات للتحضير

دعوة لمنع حلفاء مبارك من الترشح

إقرار قانون وإجراء الانتخابات بمصر

مصر تقر قانون الانتخابات

أهم أخبار الصفحة الرئيسية

المعارضة اليمينية ترفض خطاب صالح

اعتقالات بإدلب ودرعا وتعزيزات بالرسن

تزايد الضحايا ومطالبات بمحاكمة صالح

العثور على مقبرة ضحايا بوسليم

العالم في يديك

The World In Your Hands

contact : mobile@aljazeera.net

Google Translate

http://translate.google.com/translate?sl=ar&tl=en&js=n&prev=_t&hl=en&ie=UTF-8&layout=interstice

Google translate Translate

Translate from: Arabic Translate into: English

الأخبار | الشخصية | المعرفة | الاقتصاد والأعمال | الحريات وحقوق | مواقع أخرى للجريدة

Search

Member About To the Press Site Contact Contact

البث الحي

شاهد البث الحي لقناة الجزيرة

Monday, 28/10/1432 e - m, corresponding to 26/9/2011 (Updated) 1:48 pm (Makkah), 22:48 (GMT)

القائمة الرئيسية

Home:Arabic

Arabic

Reactions

Site Services

Government of Egypt Approves the amendment to election law



Young men holding up a sign describing the electoral law before you modify it as misleading (Alice)

On Sunday said the government approved the amended electoral law on the basis of the forthcoming parliamentary elections will be held before the end of the year, dominated the voting system based on individual system, after strong objections to the previous version of the law.

ملاحظات

- Proposed amendment to Egypt's elections system
- Egypt Bnovmbr ballot and meetings to prepare
- Call to prevent the allies of Mubarak's candidacy
- Adoption of the law and procedures of the elections in Egypt
- Egypt recognizes the election law

أهم أخبار الصفحة الرئيسية

- Yemeni opposition rejects speech in favor of
- Idleb arrests and shield and reinforcements

ال

أقرت الحكومة المصرية الأحد تعديل القانون الانتخابي الذي على أساسه ستجرى الانتخابات البرلمانية المرتقبة قبل نهاية العام، وغلبت فيه الاقتراع بنظام القائمة على النظام الفردي، وذلك بعد اعتراضات قوية على النسخة السابقة من القانون.

The Egyptian government approved on Sunday the electoral law amendment on the basis of which the forthcoming parliamentary elections will be held before the end of the year,

Google translate

On Sunday said the government approved the amended electoral law on the basis of the forthcoming parliamentary elections will be held before the end of the year, dominated the voting system based on individual system, after strong objections to the previous version of the law.




Lecture Plan


1. A bit more course overview [5 mins]
2. Briefly learn the history of Machine Translation [10 mins]
3. Learn about translation with one example [10 mins]
4. **Speech recognition & Applying it to MT: The noisy channel model [10 mins]**
[Stretch! **Emergency time reserves:** 5 mins]
5. Parallel-text word alignments: the IBM models [30 mins]

Speech Recognition: Acoustic Waves

- Human speech generates a wave
 - like a loudspeaker moving its magnet
- A wave for the words “speech lab” looks like:




“l” to “a”
transition:



From Simon Arnfield's web tutorial on speech, Sheffield:
<http://www.psyc.leeds.ac.uk/research/cogn/speech/tutorial/>

Acoustic Sampling

- 10 ms frame (ms = millisecond = 1/1000 second)
- ~25 ms window around frame [wide band] to allow/ smooth signal processing – it let's you see formants




...


Result:
Acoustic Feature Vectors
(after transformation,
numbers in roughly \mathbf{R}^{14})

Spectral Analysis

- Frequency gives pitch; amplitude gives volume
 - sampling at ~8 kHz phone, ~16 kHz mic (kHz=1000 cycles/sec)



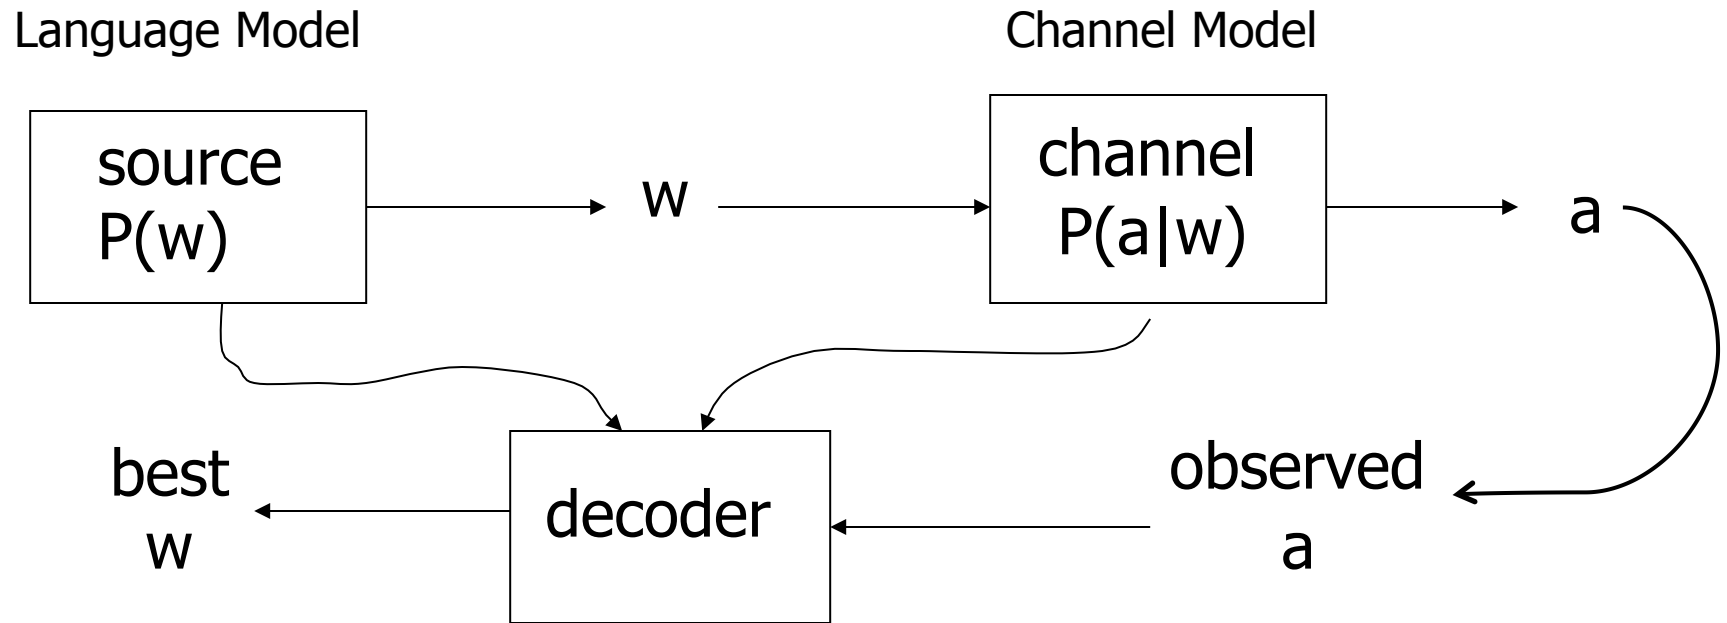
- Fourier transform of wave displayed as a spectrogram
 - darkness indicates energy at each frequency
 - hundreds to thousands of frequency samples



The Speech Recognition Problem

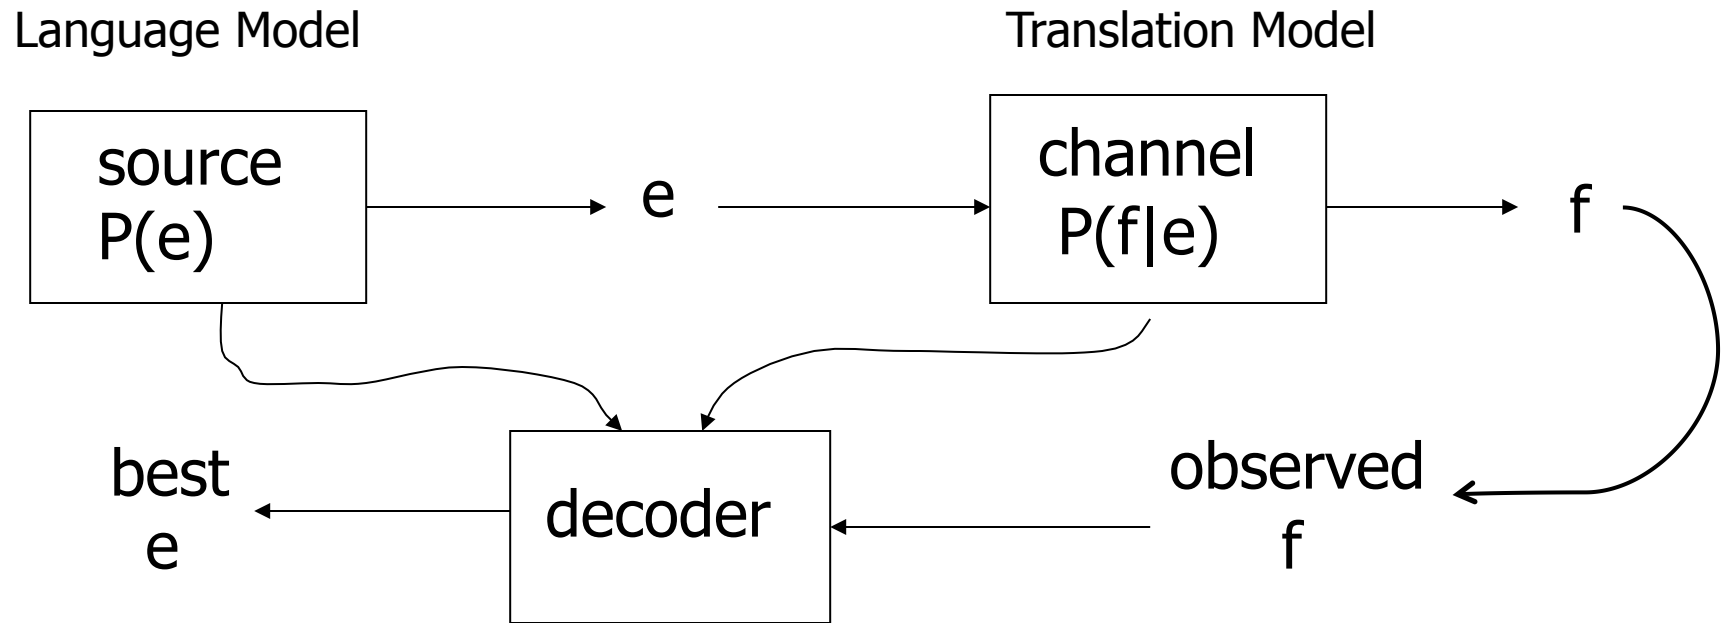
- **Speech Recognition by the Noisy Channel Model**
 - Build a generative model of encoding: We started with English words, they are transmitted as an audio signal, and we now wish to decode what we hear.
 - Listener finds most likely sequence \mathbf{w} of “words” given the sequence of acoustic observation vectors \mathbf{a}
- Use this **generative model** to decode:
- That is, use Bayes Rule

ASR System Components



$$\begin{aligned}\hat{\mathbf{w}} &= \operatorname{argmax}_{\mathbf{w}} P(\mathbf{w} | \mathbf{a}) = \operatorname{argmax}_{\mathbf{w}} \frac{P(\mathbf{a} | \mathbf{w})P(\mathbf{w})}{P(\mathbf{a})} \\ &= \operatorname{argmax}_{\mathbf{w}} P(\mathbf{a} | \mathbf{w})P(\mathbf{w})\end{aligned}$$

MT System Components



$$\begin{aligned}\hat{\mathbf{e}} &= \operatorname{argmax}_{\mathbf{e}} P(\mathbf{e} | \mathbf{f}) = \operatorname{argmax}_{\mathbf{e}} \frac{P(\mathbf{f} | \mathbf{e})P(\mathbf{e})}{P(\mathbf{f})} \\ &= \operatorname{argmax}_{\mathbf{e}} P(\mathbf{f} | \mathbf{e})P(\mathbf{e})\end{aligned}$$

Other Noisy-Channel Processes

- Handwriting recognition

$$P(\textit{text} \mid \textit{strokes}) \propto P(\textit{text})P(\textit{strokes} \mid \textit{text})$$

- OCR

$$P(\textit{text} \mid \textit{pixels}) \propto P(\textit{text})P(\textit{pixels} \mid \textit{text})$$

- Spelling Correction

$$P(\textit{text} \mid \textit{typos}) \propto P(\textit{text})P(\textit{typos} \mid \textit{text})$$

Statistical MT

Pioneered at IBM in the early 1990s



Let's make a probabilistic model of translation

$P(e | f)$

Suppose f is *de rien*

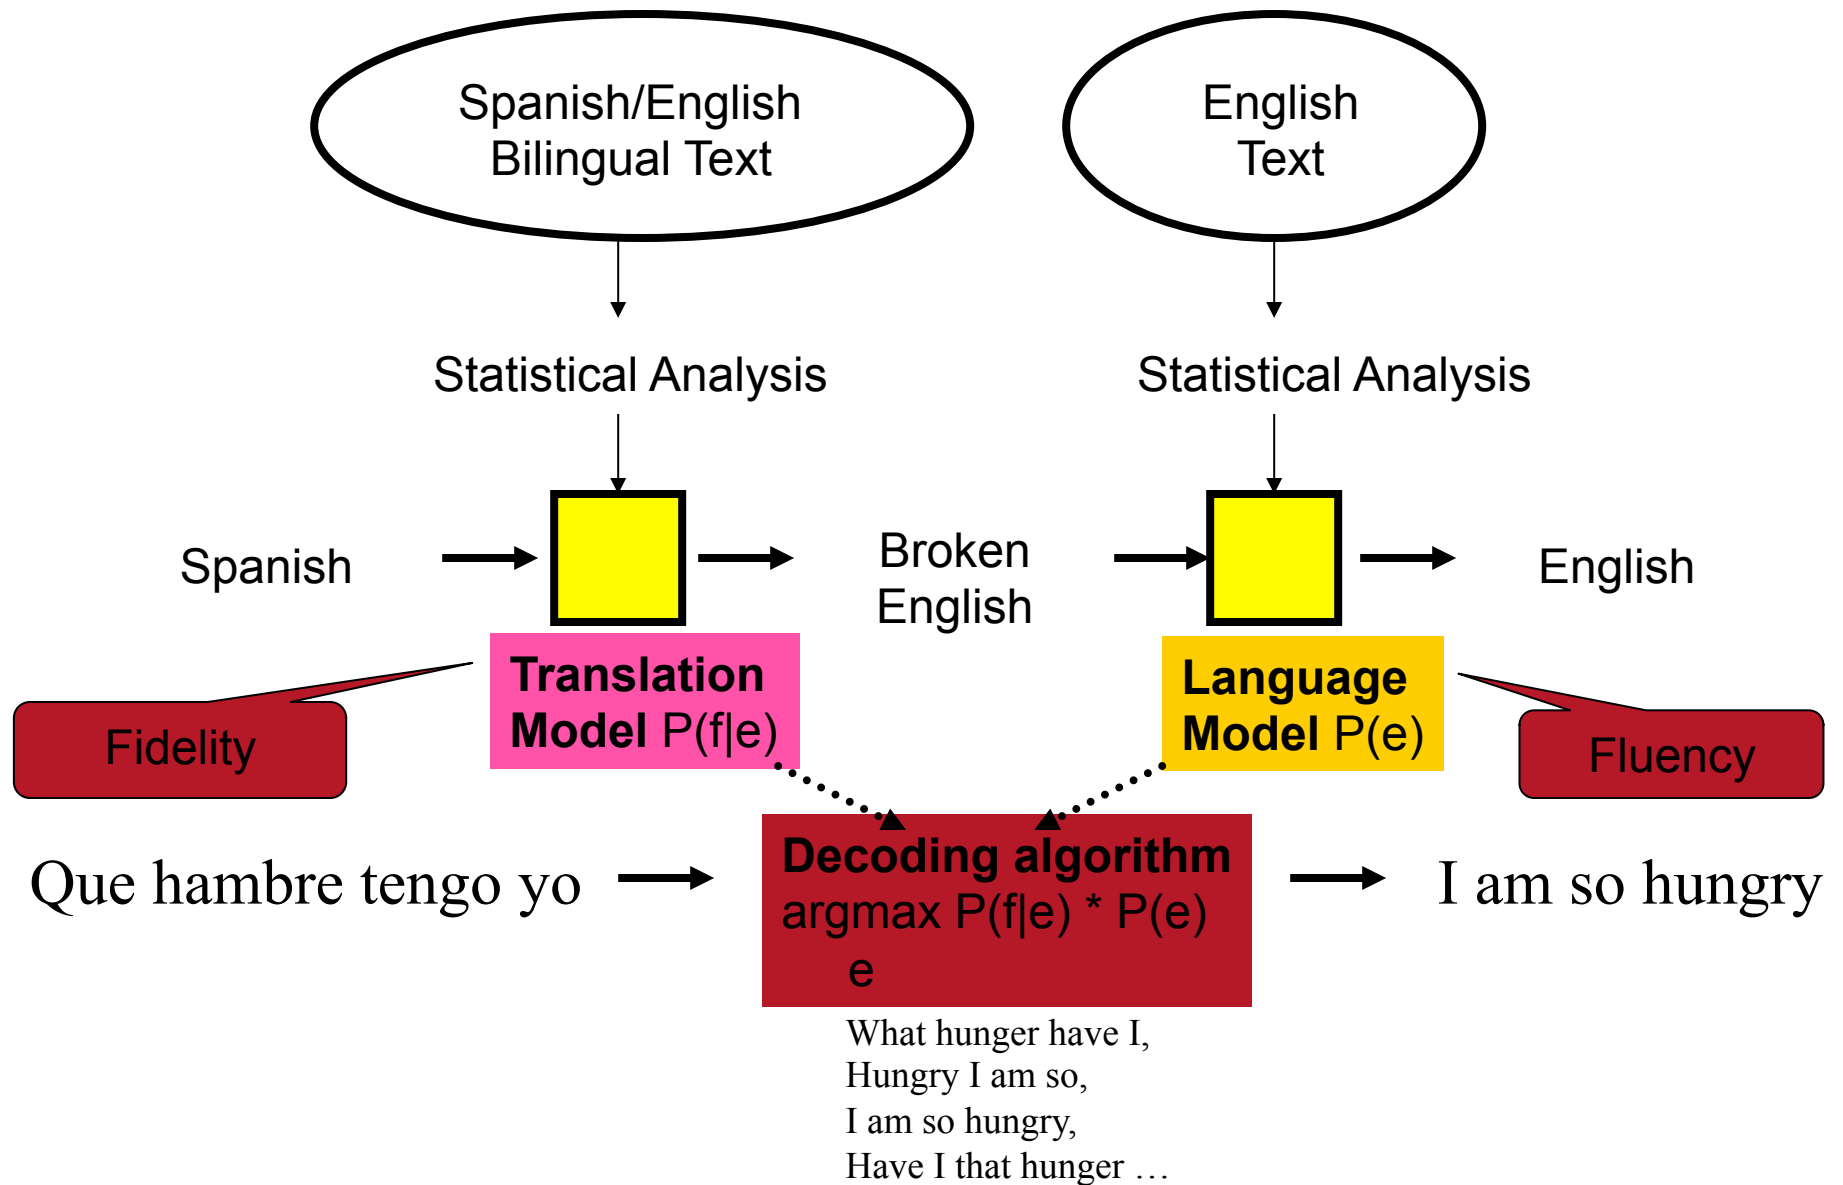
$P(\textit{you're welcome} | \textit{de rien}) = 0.45$

$P(\textit{nothing} | \textit{de rien}) = 0.13$

$P(\textit{piddling} | \textit{de rien}) = 0.01$

$P(\textit{underpants} | \textit{de rien}) = 0.000000001$

A Division of Labor



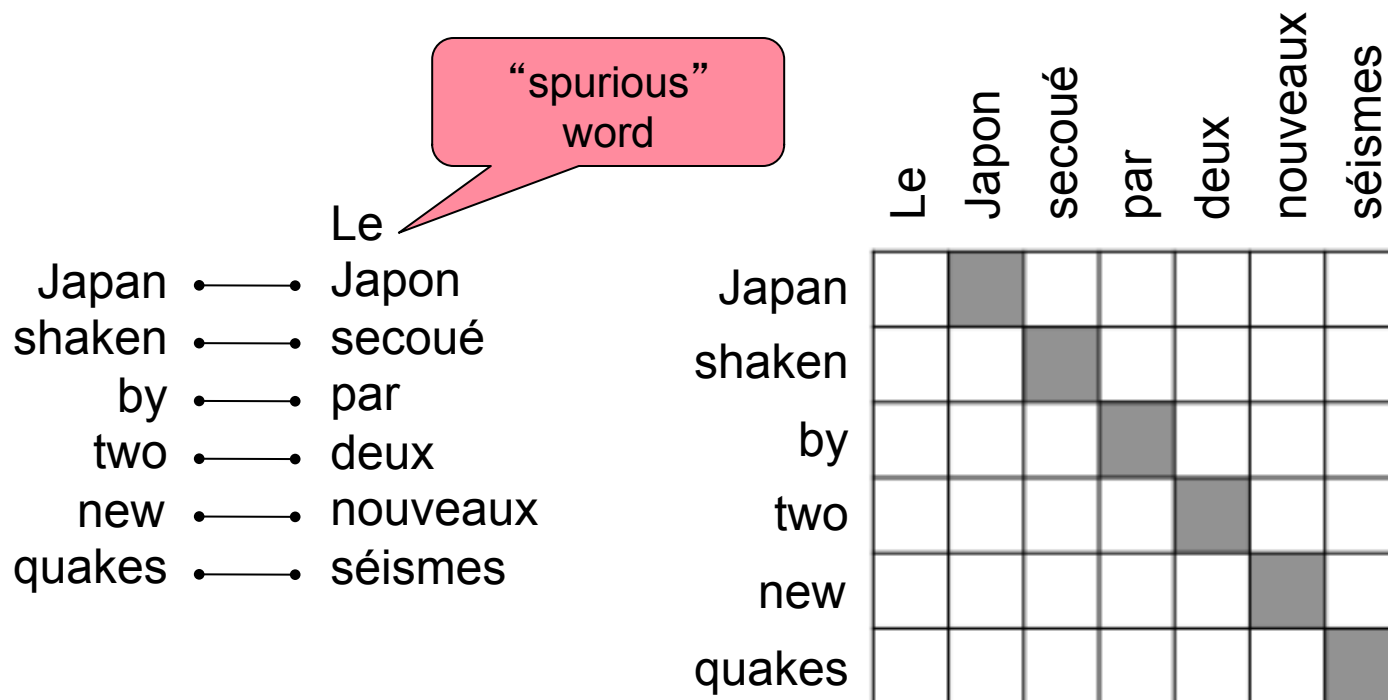


Lecture Plan

1. A bit more course overview [5 mins]
2. Briefly learn the history of Machine Translation [10 mins]
3. Learn about translation with one example [10 mins]
4. Speech recognition & Applying it to MT: The noisy channel model [10 mins]
[Stretch! **Emergency time reserves:** 5 mins]
5. **Parallel-text word alignments: the IBM models [30 mins]**

Alignments

We can factor the translation model $P(f | e)$ by identifying *alignments* (correspondences) between words in f and words in e



Alignments: harder

“zero fertility” word
not translated


And
the
program
has
been
implemented

Le
programme
a
été
mis
en
application

one-to-many
alignment

	Le	programme	a	été	mis	en	application
And							
the							
program							
has							
been							
implemented							


Alignments: harder



many-to-one
alignments

	Le	reste	appartenait	aux	autochtones
The					
balance					
was					
the					
territory					
of					
the					
aboriginal					
people					

Alignments: hardest



many-to-many
alignment

	Les	pauvres	sont	démunis
The				
poor				
don't				
have				
any				
money				

phrase
alignment

Alignment as a vector

	j		i		$a_i=j$
Mary	1	←	1	Maria	1
did	2		2	no	3
not	3	←	3	daba	4
slap	4	←	4	una	4
			5	botefada	4
			6	a	0
the	5	←	7	la	5
green	6	←	8	bruja	7
witch	7	←	9	verde	6

- used in all IBM models
- Have $f = f_1 \dots f_j \dots f_m$
- $e = e_1 \dots e_j \dots e_l$
- a is vector of length m
- maps indexes i to indexes j
- each $a_i \in \{0, 1 \dots l\}$
- e_0 is special NULL word
- $a_i = 0 \Leftrightarrow f_i$ is “spurious”
- no one-to-many alignments
- no many-to-many alignments
- but provides foundation for phrase-based alignment

IBM Model 1 generative story

	Le	programme	a	été	mis	en	application
And							
the	■						
program		■					
has			■				
been				■			
implemented					■	■	■
a_j	2	3	4	5	6	6	6

Want: $P(f|e)$ – the channel model

Given English sentence e_1, e_2, \dots, e_l

Choose length m for French sentence

For each i in 1 to m :

- Choose a_i uniformly from $0, 1, \dots, l$
- Choose f_i by translating e_{a_i}

We want to learn
how to do this

IBM Model 1 parameters

	Le	programme	a	été	mis	en	application
And							
the							
program							
has							
been							
implemented							
	a_1	2	3	4	5	6	6

$$P(f, a|e) = P(m|\ell) \prod_i P(a_i) t(f_i|e_{a_i})$$

$$= \epsilon \prod_i P(a_i) t(f_i|e_{a_i})$$

$$= \epsilon \prod_i \frac{1}{\ell + 1} t(f_i|e_{a_i})$$

$$= \frac{\epsilon}{(\ell + 1)^m} \prod_i t(f_i|e_{a_i})$$

Applying Model 1*

$P(f, a | e)$ can be used as a *translation model* or an *alignment model*

As translation model

$$P(f|e) = \sum_a P(f, a|e)$$

As alignment model


$$\begin{aligned} P(a|e, f) &= \frac{P(f, a|e)}{P(f|e)} \\ &= \frac{P(f, a|e)}{\sum_{a'} P(f, a'|e)} \end{aligned}$$

* Actually, any $P(f, a | e)$, e.g., any IBM model

Unsupervised Word Alignment

Input: a *bitext*: pairs of translated *sentences*


Output: *alignments*: pairs of translated *words*



Starting point: All word alignments equally likely


All $P(\text{french-word} \mid \text{english-word})$ equally likely

Unsupervised Word Alignment



"la" and "the" observed to co-occur frequently,
so $P(\text{la} \mid \text{the})$ is increased.

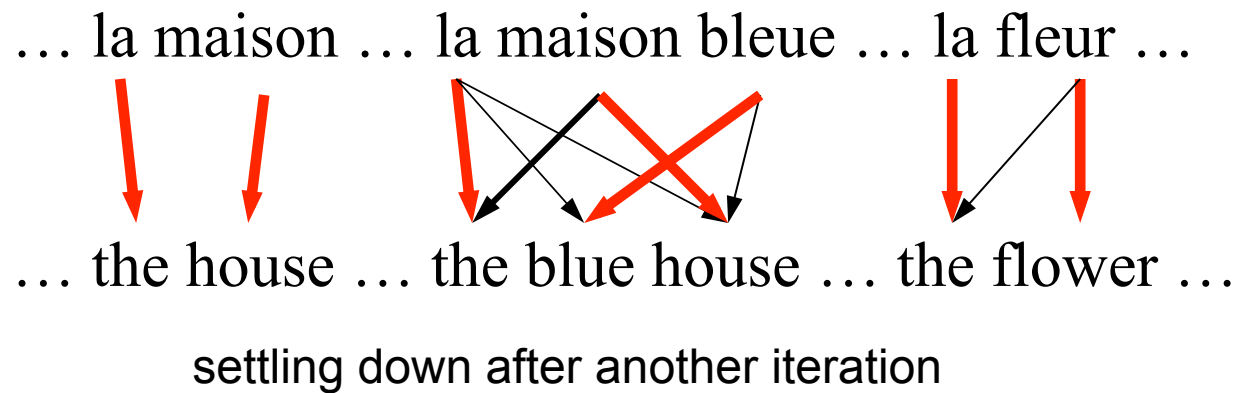
Unsupervised Word Alignment



“maison” co-occurs with both “the” and “house”, but $P(\text{maison} \mid \text{house})$ can be raised without limit, to 1.0, while $P(\text{maison} \mid \text{the})$ is limited (see 3rd example)

(pigeonhole principle)

Unsupervised Word Alignment



That was the idea of IBM Model 1 !

Model 1: Word alignment learning with Expectation-Maximization (EM)

- Start with $t(f^p|e^q)$ uniform, including $P(f^p|\text{NULL})$
- For each sentence pair (e, f)
 - For each French position i

- Calculate posterior over English positions $P(a_i | e, f)$

$$P(a_i = j | f, e) = \frac{t(f_i | e_j)}{\sum_{j'} t(f_i | e_{j'})}$$

- Increment count of word f_i translating each word e_{a_i}
 - $C(f_i | e_j) += P(a_i = j | f, e)$


- Renormalize counts to give probs $t(f^p|e^q) = \frac{C(f^p|e^q)}{\sum_{f^x} C(f^x|e^q)}$
- Iterate until convergence

IBM Models 1,2,3,4,5

- Models for $P(\mathbf{f}|\mathbf{e})$ and $P(\mathbf{a}|\mathbf{f},\mathbf{e})$ via $P(\mathbf{f},\mathbf{a}|\mathbf{e})$
- There is a set of English words and the extra English word NULL
- Each English word generates and places 0 or more French words
- Any remaining French words are deemed to have been produced by NULL
 - “Spurious” words

IBM Models 1,2,3,4,5

- In Model 2, the placement of a word in the French depends on where it was in the English



- Unlike Model 1, Model 2 captures the intuition that translations should usually “lie along the diagonal”

- A main focus of PA #1

IBM Models 1,2,3,4,5

- In Model 3, we model how many French words an English word can produce, using a concept called *fertility*


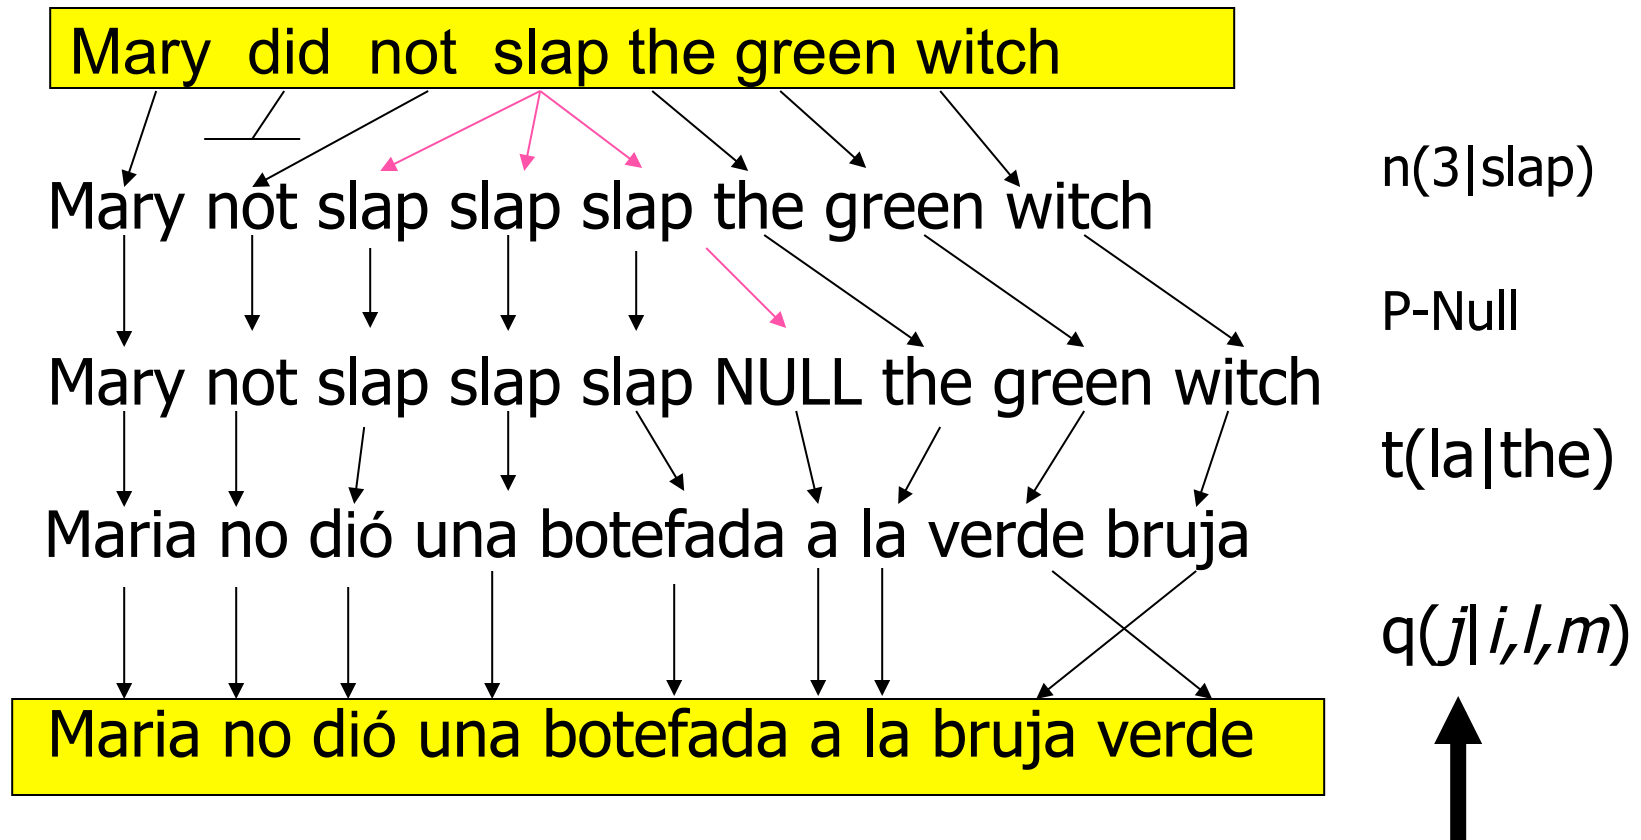


Figure 32.3

Alignment example.

Model 3 generative story



Probabilities can be learned from raw bilingual text.

IBM Model 3 (from Knight 1999)

- For each word e_j in English sentence, choose a **fertility** ϕ_j . The choice of ϕ_j depends only on e_j , not other words or ϕ 's: $n(\phi_j | e_j)$
- For each word e_j , generate ϕ_j French words. Choice of French word depends only on English word e_j , not on English context or any other French words.
- Permute all the French words. Each French word gets assigned absolute target position slot (1,2,3, etc.). Choice of French word position dependent only on absolute position of English word generating it and sentence lengths

Model 3: $P(f|e)$ parameters

- What are the parameters for this model?
- **Word translation**: $t(\text{casa} | \text{house})$
- **Spurious words**: $t(f_i | \text{NULL})$
- **Fertilities**: $n(1|\text{house})$: prob that “house” will produce 1 Spanish word whenever it appears.
- **Distortions**: $q(5|2,4,6)$: prob that word in position 2 of French translation was generated by word in position 5 of English sentence, given that 4 is length of English sentence, 6 is French length

Spurious words

- We could have $n(3|\text{NULL})$ (probability of there being exactly 3 spurious words in a French translation)
 - But seems wrong...
- Instead, of $n(0|\text{NULL})$, $n(1|\text{NULL}) \dots n(25|\text{NULL})$, have a single parameter p_1
- After assign fertilities to non-NULL English words we want to generate (say) z French words.
- As we generate each of z words, we optionally toss in spurious French word with probability p_1
- Probability of not adding spurious word: $p_0 = 1 - p_1$

Distortion probabilities for spurious words

- Shouldn't just have $q(0|5,4,6)$, i.e., chance that source position for word 5 is position 0 (NULL).
- Why? These are spurious words! Could occur anywhere!! Too hard to predict
- Instead,
 - Use normal-word distortion parameters to choose positions for normally-generated French words
 - Put NULL-generated words into empty slots left over
 - If three NULL-generated words, and three empty slots, then there are $3!$, or six, ways for slotting them all in
 - We'll assign a probability of $1/6$ for each way!

Model 3 parameters

- n, t, p, q
- Again, if we had complete data of English strings and step-by-step rewritings into Spanish, we could:
 - Compute $n(0|did)$ by locating every instance of “did”, and seeing how many words it translates to
 - $t(maison|house)$ how many of all French words generated by “house” were “maison”
 - $q(5|2,4,6)$ out of all times some second word is in a translation, how many times did it come from the fifth word (in sentences of length 4 and 6 respectively)?

Since we don't have word-aligned data...

- We bootstrap alignments from incomplete data
- From a sentence-aligned bilingual corpus
 - 1) Assume some startup values for n , q , t , p .
 - 2) Use values for n , q , t , p in model 3 to work out chances of different possible alignments. Use these alignments to update values of n , q , t , p .
 - 3) Go to 2
- This is a more complicated case of the EM algorithm

Difficulty: Alignments are no longer independent of each other. Have to use approximate inference

Examples: translation & fertility

the

f	$t(f e)$	ϕ	$n(\phi e)$
le	0.497	1	0.746
la	0.207	0	0.254
les	0.155		
l'	0.086		
ce	0.018		
cette	0.011		


not

f	$t(f e)$	ϕ	$n(\phi e)$
ne	0.497	2	0.735
pas	0.442	0	0.154
non	0.029	1	0.107
rien	0.011		

farmers

f	$t(f e)$	ϕ	$n(\phi e)$
agriculteurs	0.442	2	0.731
les	0.418	1	0.228
cultivateurs	0.046	0	0.039
producteurs	0.021		

Example: idioms

he is noding

 il hoche la tête

nodding

f	$t(f e)$	ϕ	$n(\phi e)$
signe	0.164	4	0.342
la	0.123	3	0.293
tête	0.097	2	0.167
oui	0.086	1	0.163
fait	0.073	0	0.023
que	0.073		
hoche	0.054		
hocher	0.048		
faire	0.030		
me	0.024		
approuve	0.019		
qui	0.019		
un	0.012		
faites	0.011		

Example: morphology

should

f	$t(f e)$	ϕ	$n(\phi e)$
devrait	0.330	1	0.649
devraient	0.123	0	0.336
devrions	0.109	2	0.014
faudrait	0.073		
faut	0.058		
doit	0.058		
aurait	0.041		
doivent	0.024		
devons	0.017		
devrais	0.013		

IBM Models 1,2,3,4,5


- In model 4 the placement of later French words produced by an English word depends on what happened to earlier French words generated by that same English word

Alignments: linguistics

On Tuesday Nov. 4, earthquakes rocked Japan once again



Des tremblements de terre ont à nouveau touché le Japon mardi 4 novembre



IBM Models 1,2,3,4,5

- In model 5 they patch model 4. They make it do non-deficient alignment. That is, you can't put probability mass on impossible things.

IBM StatMT Translation Models

- IBM1 – lexical probabilities only
 - IBM2 – lexicon plus absolute position
 - HMM – lexicon plus relative position
 - IBM3 – plus fertilities
 - IBM4 – inverted relative position alignment
 - IBM5 – non-deficient version of model 4
-
- All the models above handle 0:1, 1:0, 1:1, 1:n alignments *only*

[Brown et al. 93, Vogel et al. 96]

Why all the models?

- We don't start with aligned text, so we have to get initial alignments from somewhere.
- The alignment space has many local maxima
- Model 1 is words only, a simple model that is relatively easy and fast to train.
- The output of M1 can be a good place to start M2
 - “Starting small”. Also, it's convex!
- The sequence of models allows a better model to be found, faster
 - The intuition is like deterministic annealing